

C-Bounded Powersort's time complexity and correctness analysis

Théo Araújo Magalhães, Victor Almeida Campos

April 5, 2025

1 Definitions:

Clarification: Definitions with a * at the end come from Munro and Wild's papers. This is a 'to-do', I will properly cite them later.

Runs: Contiguous regions of the input that are already in sorted order. Let $R_0, \dots, R_{r-1}$ be the runs in the input, L_i denotes the length of R_i .*

Run-adaptive mergesort: Any stable run-adaptive mergesort can be described as a merge tree T : an (ordered rooted) tree with leaves $R_0, \dots, R_{r-1}$ (appearing in that order from left to right). Each internal node v of T corresponds to the result of merging its children.*

Merge costs: $M(v)$ is the total length of the resulting merge corresponding to the internal node v . The total merge cost of T is then:

$$M(T) = \sum_{v \in T} M(v). \quad (1)$$

Run boundaries: Denoted by B_i , it is the boundary between the $(i - 1)$ st and i th run ($i=1, \dots, r - 1$). In binary merge trees, there is a one-to-one correspondence between the B_i and the (inner/non-leaf) merge-tree nodes.*

Node depth: Let d_i be the depth of leaf R_i in T , where depth is the number of edges on the path to the root. Every element in run R_i is counted exactly d_i times in $\sum_v M(v)$, so:

$$M = \sum_{i=0}^{r-1} d_i \cdot L_i \quad (2)$$

Node power: Let $\alpha_0, \dots, \alpha_m, \sum \alpha_j = 1$ be leaf probabilities. For $1 \leq j \leq m$, let j be the internal node separating the $(j - 1)$ st and j th leaf. The power of (the split at) node j is:

$$P_j = \min\{\ell \in \mathbb{N} : \lfloor \frac{a}{2^{-\ell}} \rfloor < \lfloor \frac{b}{2^{-\ell}} \rfloor\}, \text{ where } a = \sum_{i=0}^{j-1} \alpha_i - \frac{1}{2} \alpha_{j-1}, b = \sum_{i=0}^{j-1} \alpha_i + \frac{1}{2} \alpha_j. \quad (3)$$

(P_j is the index of the first bit where the (binary) fractional parts of a and b differ.)

Powersort: The intuition behind Powersort is to imagine a complete binary tree T_v sitting on top of the array, where each element $A[i]$ of the input corresponds to

a leaf of T_v . T_v corresponds to the merge tree of a non-adaptive two-way Mergesort (starting with individual elements). Now each run boundary B_j also corresponds to an (inner) node w of this (virtual) k -ary tree T_v , namely the lowest common ancestor of the two leafs adjacent to B_j . *[We need to re-explain this part better so that we can talk about the concept of a “Powersort merge tree” or something of the sorts. Helps with the proof of our main proposition]

Forgetful K-Stack: A forgetful stack is a stack that forgets its k -th element if it already has k elements in it and a new one is added.

C-Bounded Powersort: The Powersort algorithm using a Forgetful C-Stack instead of the traditional stack with size $\log n + 1$.


Processing a subtree of height h : Let T be a merge tree, T' one of its subtrees of size h and v the root of T' . Processing T' is executing all the merges necessary to fill node v , given that it has length $M(v)$.

2 Proofs:

Proposition: Let T be a Powersort merge tree of height H . One pass of the C-Bounded Powersort algorithm is enough to process all subtrees of T with height of $C-1$.

Proof by induction:

Base case - T of height 2:



Let v be the root of T and v_1 and v_2 its two children. Therefore, T has two subtrees, T' and T'' , both of height one. For the left subtree T' , v_1 is the root and it has only two leaf children: the runs R_1 and R_2 . Given that $h = 1$, a Forgetful 2-Stack will work with 2-Bounded Powersort as follows:

R_1 will be found and added to the stack, then, given that there is nothing to compare R_1 with, the algorithm will search for the adjacent run, R_2 . After finding it, the algorithm will calculate the power between R_1 and R_2 , which is equal to the height of v_1 . Given that we only have one power calculated, this will lead to no merges and R_2 will be added to our stack. Afterwards, the algorithm will proceed to the right T'' subtree, which is rooted in v_2 . The algorithm will find R_3 and, then, calculate the power between R_2 and R_3 . This will result in the height of v , which is smaller than that of v_1 , our previously calculated power, leading to the merge of runs R_1 and R_2 into run R_{12} .

After this, our Forgetful 2-Stack will contain R_{12} and R_3 . Then, the algorithm will find run R_4 . Calculating the power between R_3 and R_4 leads to height v_2 , which is bigger than that of v , our previously calculated power, meaning the algorithm will not merge any runs and will add R_4 to the Forgetful 2-Stack. This will cause the stack to “forget” about R_{12} , meaning it will only contain R_3 and R_4 . Once we find R_4 , we will have reached the end of the array.

Following the logic of the algorithm, once the scan reaches the end of the array everything that still is in the Forgetful-Stack will be merged, thus, in this situation, a merge between R_3 and R_4 will happen, leading to the processing of both subtrees T' and T'' . This concludes our base case.

Hypothesis: Suppose the proposition holds true for all subtrees of height $h \leq k$ in a powersort merge tree.

Inductive Step: Let T be a Powersort merge tree of height H . We will prove that a $(k+2)$ -Bounded Powersort processes all of T 's subtrees that have height of $h = k+1$. Let T' be one of those subtrees and v its root. We will analyze the sub-trees T_ℓ and T_r , each rooted in the left and right children of v , v_ℓ being the left child and v_r the right child. The height of the T_ℓ is, at most, k , therefore, we can apply our induction hypothesis and process T_ℓ , thus, v_ℓ will contain run R_ℓ , which is the result of the merge of all runs in T_ℓ . The processing of this sub-tree will happen once the most-left run of T_r is analyzed by the algorithm.¹ By the end of the processing of T_ℓ , our Forgetful $(k+2)$ -Stack will contain R_ℓ and the most-left run of T_r . Given that, by hypothesis, the processing of T_r requires one pass of a $(k+1)$ -Bounded Powersort^{2 3}, and that our Forgetful $(k+2)$ -Stack has exactly k spaces left, with one of them already being occupied by one run of T_r , the algorithm will be able to process T_r . Similarly to what happened during the processing of T_ℓ , the calculation of the power between the right-most run of T_r (which is also the right-most run of T') and the next run, which is the left-most run of another subtree of T with height $k+1$, will result in the height of $v - 1$.⁴ Given that all the powers calculated between runs in T' are greater than $v-1$, the algorithm will merge all the remaining runs in our Forgetful $(k+2)$ stack, leading to the processing of T_r but also to the processing of T' , given that R_ℓ will still be on the stack alongside all runs from T_r . Therefore, if the $(k+2)$ -Bounded Powersort was able to process a generic T' subtree of T with height $k+1$, it is fair to assume that it will be able to further process all the remaining subtrees of same height by following the same logic as the one described for T' . $\square$

¹Once the power between the most-right run of T_ℓ with most-left run of T_r is calculated, it will result in the height of v , which is smaller than the height of v_ℓ , meaning the algorithm will merge everything left in our Forgetful $k+2$ -Stack, for all powers calculated between runs in T_ℓ is greater than v_ℓ

²Remember that T_ℓ and T_r have, at max, height of k

³Requiring a $(k+1)$ -Bounded Powersort means the processing needs $k+1$ free spaces in the Forgetful C-Stack, not a specific tuning of the algorithm.

⁴This is due to the fact that both runs are in adjacent subtrees: v is the left child of a node v_p and there is a node v' which is the right child of the same v_p , the first common ancestor between v and v' is v_p , therefore, the first common ancestor between runs in each subtree (the subtree rooted in v and the subtree rooted in v') is v_p

Proposition: It takes C-Bounded Powersort at most $\mathcal{O}(n \cdot \log n)$ element scans to order an array of size n .

Proof. Each pass through the array scans element i , with $0 \leq i \leq n$, at least once. Given that the algorithm will do $\frac{\log n}{c-1}$ passes⁵, i will be read at least $\frac{\log n}{c-1}$ times. Once the entire array is ordered, i will also have been read at least twice for each merge it was present in, i.e, $2 \cdot d_i$, with d_i being equal to the height of T . Therefore, the total number of scans is:

$$Scans \leq \sum_{i=0}^n 2 \cdot \log n + n \cdot \frac{\log n}{c-1} \quad (4)$$

$$Scans \leq n \log n \cdot (2 + \frac{1}{c-1}) \quad (5)$$

$$Scans \leq \mathcal{O}(n \cdot \log n) \quad (6)$$

□

⁵The height of a Powersort merge tree T is $\log n$ and the C-Bounded Powersort algorithm starts by processing subtrees of T that have height less than or equal to $c-1$. This results in a smaller tree T' of height $\log n - (c-1)$. A second pass will, once again, process the subtrees of T' that have height less than or equal to $c-1$, resulting in a smaller tree T'' . Therefore, if k is the amount of trees created by this process until the last tree is of height less than or equal to $c-1$, meaning it will only require one last pass of the algorithm to be processed, $k \cdot (c-1) \leq \log n$

3 Pseudocodes:

TO-DOs:

1. Os comentários para cada algoritmo precisam estar de acordo com as definições usadas no primeiro esboço do artigo (powersort.tex) e precisam ser revisados para fazerem sentido, alguns ainda estão bem abstratos e/ou não descrevem corretamente os algoritmos.
2. Tanto para esse pdf quanto para o anterior, ou seja, para o artigo como um todo, o nome do professor Victor se enquadra como autor? Imagino que sim. Essa pergunta é meio boba.
3. Código para a estrutura de dados que iremos usar.
4. Reestruturar o vai e volta.

Planejamento da estrutura: Guardaremos o início da primeira run, seu fim e seu power com a run anterior (se houver), além disso, cada elemento seguinte será: o fim da run e o power da run com a anterior. run e não teremos um elemento sentinela. O início de uma run baseia-se no fim da anterior + 1. Assim que incluímos um elemento que estoura o tamanho da pilha, temos que esquecer o primeiro elemento e ajustar o novo primeiro elemento para também guardar o seu início. (cada pop ou insert decrementa ou aumenta o tamanho)

Como vamos fazer scans da esquerda para a direita e da direita para a esquerda, precisamos alternar entre somar um no fim de uma run para achar o começo da próxima e diminuir um. ← pensar nisso

Comentário do professor: “Só um detalhe de notação, usa aqui (A, s_2, n) . $A[s_2]$ é o valor e não o índice.”

Resposta: Fiz desse jeito porque é a forma que eles escreveram no artigo, mas achei bem estranho também, depois vou mudar.

Introduction: We are very smart and we know things, this is the complaints we have for things that weren't done by us and how we would have made them better.

3.1 Powersort Versions:

3.1.1 Standard Powersort

Commentary on Wild and Munro's version: A single pass algorithm that sorts the entire array.

Algorithm 1: Powersort($A[1..n]$)	
1:	$X :=$ stack of runs; //capacity $\lceil \lg n(n) \rceil + 1$
2:	$P :=$ stack of powers; //capacity $\lceil \lg n(n) \rceil + 1$
3:	$s_1 \leftarrow 1; e_1 \leftarrow \text{ExtendRunRight}(A[1], n);$
4:	while $e_1 < n$ do
5:	$s_2 \leftarrow e_1 + 1; e_2 \leftarrow \text{ExtendRunRight}(A[s_2], n);$ //$A[s_2..e_2]$ next run
6:	$p \leftarrow \text{NodePower}(s_1, e_1, s_2, e_2, n);$ //P_j for node j between $A[s_1..e_1]$
	and $A[s_2..e_2]$
7:	while $P_{\text{top}}() > p$ do
8:	//previous merge deeper in tree than current
9:	$P_{\text{pop}}();$ //$\rightarrow$ merge and replace run $A[s_1..e_1]$ by result
10:	$(s_1, e_1) \leftarrow \text{Merge}(X.\text{pop}(), A[s_1..e_1]);$
11:	$X.\text{push}(A[s_1, e_1]);$
12:	$P.\text{push}(p);$
13:	$s_1 \leftarrow s_2; e_1 \leftarrow e_2;$ //Now $A[s_1...e_1]$ is the rightmost run
14:	while $\neg X.\text{empty}()$ do
15:	$(s_1, e_1) \leftarrow \text{Merge}(X.\text{pop}(), A[s_1..e_1]);$

3.1.2 First variation of C-Bounded Powersort (only goes right)

Commentary on our first version: This algorithm performs multiple passes on the array, processing all subtrees of height $C-1$ with each pass. We remind the reader of how Powersort works: It is implicitly a complete binary tree of size $\log n$. Therefore, the max amount of passes needed to sort the array is:

$$\text{N}^\circ \text{ of scans} \leq \frac{\log n}{C-1} \quad (7)$$

Algorithm 2: First_C-Bounded_Powersort($A[1..n]$)	
1:	$X :=$ C-bounded stack of runs; //capacity C
2:	$P :=$ C-bounded stack of powers; //capacity C
3:	$s_1 \leftarrow 1; e_1 \leftarrow \text{ExtendRunRight}(A[1], n);$
4:	while $e_1 < n$ do
5:	//The while above can be tweaked to use max amount of passes
6:	while $e_1 < n$ do
7:	//We do a pass covering the entire array
8:	$s_2 \leftarrow e_1 + 1; e_2 \leftarrow \text{ExtendRunRight}(A[s_2], n);$
9:	$p \leftarrow \text{NodePower}(s_1, s_2);$ //Using our node power algorithm
10:	while $P_{\text{top}}() > p$ do
11:	//We follow the same merge policy as Munro and Wilde
12:	$P_{\text{pop}}();$
13:	$(s_1, e_1) := \text{Merge}(X_{\text{pop}}(), A[s_1..e_1]);$
14:	$X_{\text{push}}(A[s_1, e_1]); P_{\text{push}}(p);$
15:	$s_1 \leftarrow s_2; e_1 \leftarrow e_2;$
16:	while $\neg X_{\text{empty}}()$ do
17:	$(s_1, e_1) := \text{Merge}(X_{\text{pop}}(), A[s_1..e_1]);$
18:	//We're using a limited stack, so we need to check for another pass
19:	$s_1 \leftarrow 1; e_1 \leftarrow \text{ExtendRunRight}(A[1], n);$ //If $e_1 = \text{end}$, the array is sorted

3.1.3 Second variation of C-Bounded Powersort (goes left and right!)

Commentary on our second version: This works rather similarly to our first version. The main difference is the fact that, once we reach the end of the array, instead of restarting our scan from the beginning, we scan, starting from the end, to the left, until we reach the beginning of the array, and, thus, we repeat the scan-right process. This is slightly better than the first version, given that it takes advantage of runs still being in cache instead of restarting from the beginning.

The code for the entire algorithm is too big to fit in a single page, therefore, we have subdivided it in three parts, as shown below.

Algorithm 3: C-Bounded_Powersort(A[1...n]) - A

```
1:  $X :=$  C-bounded stack of runs and powers;
2:  $s_0 :=$  This is a temp variable for the merges;
3:  $s_1 \leftarrow 1; e_1 \leftarrow \text{ExtendRunRight}(1, n);$ 
4: while true do
5:   while  $e_1 < n$  do
6:      $s_2 \leftarrow e_1 + 1; e_2 \leftarrow \text{ExtendRunRight}(s_2, n);$ 
7:     //We do not need variable  $s_2$ , but it helps with clarity.
8:      $p \leftarrow \text{NodePower}(s_1, s_2);$  //  $s_2$  is just  $e_1 + 1$ 
9:     while  $X.size > 1$  and  $X.top().power > p$  do
10:       $s_0 \leftarrow X.pop().boundary;$ 
11:      Merge(A,  $s_0, s_1, e_1$ );
12:       $s_1 \leftarrow s_0;$ 
13:      X.push( $s_1, p$ );
14:       $s_1 \leftarrow s_2; e_1 \leftarrow e_2;$ 
15:   while  $X.size > 1$  do
16:      $s_0 \leftarrow X.pop().boundary;$ 
17:     Merge(A,  $s_0, s_1, e_1$ );
18:      $s_1 \leftarrow s_0;$ 
19:    $s_1 \leftarrow X.pop().boundary;$ 
20:    $s_0 \leftarrow s_1; s_1 \leftarrow e_1; e_1 \leftarrow s_0;$ 
21:   //The merge of all runs left in the stack results in a
   single run at the end of the array. We use this run as
   the first run in the left pass. This saves on doing a
   call for scanleft to 'switch sides'.
22:   if  $e_1 == 1$  then
23:     break;
24:     //If the run resulting from the merges reaches the
     beginning of the array, it is the entire array.
25:   while  $e_1 > 1$  do
26:     //We're scanning from n down to 1!
27:      $s_2 := e_1 - 1; e_2 := \text{ExtendRunLeft}(s_2, 1);$ 
28:      $p := \text{NodePower}(s_1, s_2);$ 
29:     while  $X.size > 1$  and  $X.top() > p$  do
30:        $s_0 \leftarrow X.pop().boundary;$ 
31:       Merge(A,  $s_0, s_1, e_1$ );
32:        $s_1 \leftarrow s_0;$ 
33:       X.push( $s_1, p$ );
34:        $s_1 \leftarrow s_2; e_1 \leftarrow e_2;$ 
35:   while  $X.size > 1$  do
36:      $s_0 \leftarrow X.pop().boundary;$ 
37:     Merge(A,  $s_0, s_1, e_1$ );
38:      $s_1 \leftarrow s_0;$ 
39:    $s_1 \leftarrow X.pop().boundary;$ 
40:    $s_0 \leftarrow s_1; s_1 \leftarrow e_1; e_1 \leftarrow s_0;$ 
41:   //Swapping for to the scan right. The equivalent of what
   we did before scanning left.
42:   if  $e_1 == n$  then
43:     break;
```

3.2 NodePower versions:

3.2.1 Standard Nodepower:

Commentary on Wild and Munro's Node Power: This is so silly and expensive!

Algorithm 4: Nodepower(s_1, e_1, s_2, e_2, n)
1: $n_1 := e_1 - s_1 + 1$; $n_2 := e_2 - s_2 + 1$; $\ell := 0$; 2: $a := (s_1 + n_1/2 - 1)/n$; $b := (s_2 + n_2/2 - 1)/n$; 3: while $\lfloor a \cdot 2^\ell \rfloor \neq \lfloor b \cdot 2^\ell \rfloor$ do $\ell := \ell + 1$; 4: return (ℓ)

3.2.2 Our Nodepower:

Commentary on our Node Power: Ours is cool and epic. It could also be reduced to a single line: `return (__builtin_clz(s1 ^ s2))` (TO-DO: explanation of why it works)

Algorithm 5: Our_Nodepower(s_1, s_2)
1: $\text{xor} \leftarrow s_1 \wedge s_2$; 2: $\ell \leftarrow _\text{builtin_clz}(\text{xor})$; 3: return (ℓ);

3.3 The data structure:

Introduction: Wild and Munro's paper shows no specific code for their stack data structure, this is probably due to the fact that they simply use a standard stack of height $\log n$. Thus, we will have no 'counterpart' algorithm to analyze. Our algorithm implements a modified version of a stack which we call:

C-Bounded Stack"/"Forgetful Stack.

3.3.1 C-Bounded Stack:

Commentary on our data structure: When creating this data structure, we had to think of two particularities:

1. This is a stack of limited size.
2. This is a stack that has to consider scans going right and a scans going left.

In order to deal with the limited size, which is the 'forgetful' aspect of the data structure, we simply use circularity: we keep the index of where the next new element will be inserted in the case of a push and increment it. If we reach the end of the array, we assign the free space to be the first element and overwrite it in case of a push. As the stack fills, we also increment a variable for 'size', which corresponds to the amount of elements in the stack. Thus, we keep two variables, one for 'top' and one for 'size'. Other than that, each element in the stack consists of a run boundary and the power of this run with its predecessor.

Run boundaries: When the algorithm is scanning right, the boundary it keeps is the start of the run found. When the algorithm is scanning left, the boundary it keeps is the end of the run found.

Algorithm 6: C-Bounded Stack(C)

```
1: top  $\leftarrow$  1; size  $\leftarrow$  0; //Starting values for the stack
2: MaxSize  $\leftarrow$  C; //Max stack size in number of elements
3: array[MaxSize]; //The array will store tuples of (index,power)
4: func pop():
5:   if size > 1: then
6:     if top == 1: then
7:       top  $\leftarrow$  MaxSize; size--;
8:       return array[top];
9:     else
10:      top--; size--;
11:      return array[top];
12: func push(boundaryi, k):
13:   //boundaryi is the index of the start of the i-th run
14:   //k is the power of the i-th run with the (i+1)-th run.
15:   array[top]  $\leftarrow$  (boundary1, k);
16:   top  $\leftarrow$  top%MaxSize + 1;
17:   if (size  $\neq$  MaxSize) then
18:     size++
19: func func Top():
20:   if size  $\geq$  1: then
21:     return array[(top == 1) ? MaxSize : (top - 1)]
22:   else
23:     return null; //This will never happen. Top is only called
                //in the algorithm if size > 1.
```